



Transformers

Slido: <https://app.sli.do/event/k5TBSCako9oeaZZZYuWkfK>

Yu Meng
University of Virginia
yumeng5@virginia.edu

Sept 29, 2025

Overview of Course Contents

- Week 1: Logistics & Overview
- Week 2: N-gram Language Models
- Week 3: Word Senses, Semantics & Classic Word Representations
- Week 4: Word Embeddings
- Week 5: Sequence Modeling & Recurrent Neural Networks (RNNs)
- **Week 6: Language Modeling with Transformers**
- Week 9: Large Language Models (LLMs) & In-context Learning
- Week 10: Knowledge in LLMs and Retrieval-Augmented Generation (RAG)
- Week 11: LLM Alignment
- Week 12: Reinforcement Learning for LLM Post-Training
- Week 13: LLM Agents + Course Summary
- Week 15 (after Thanksgiving): Project Presentations

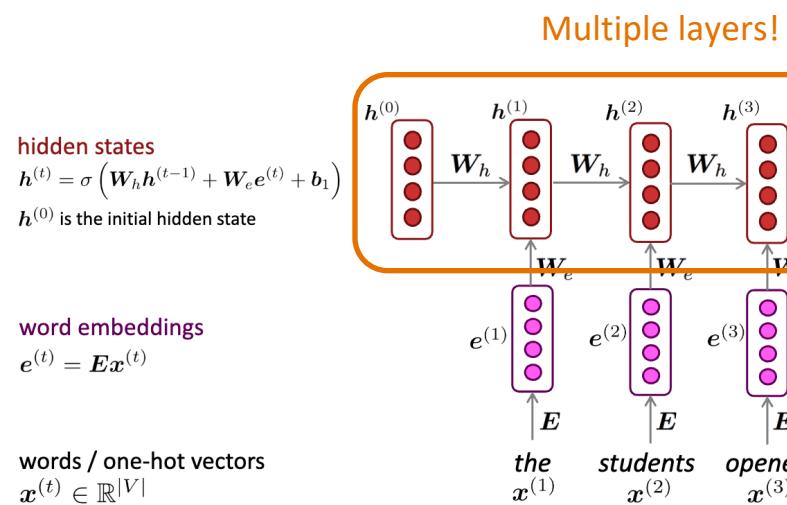
Reminder

- Project proposal grades & feedback released
- Assignment 3 released; due date: 10/06 11:59pm

(Recap) Sequence Modeling: Overview

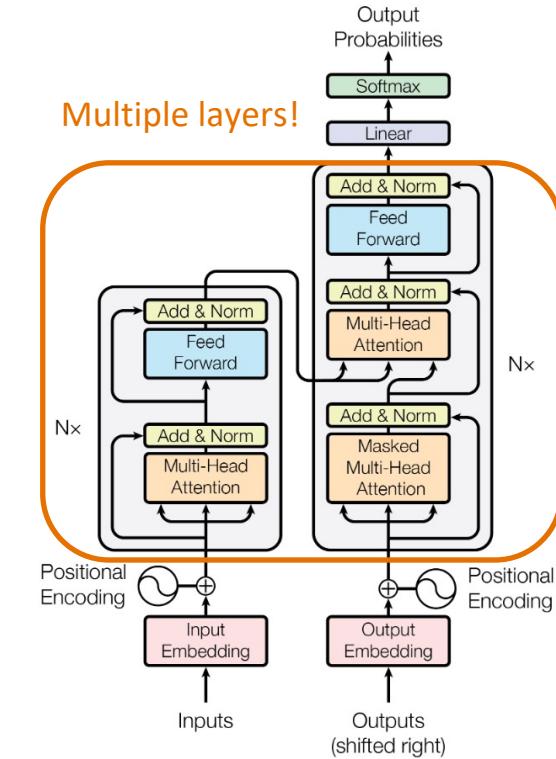
- Use deep learning methods to understand, process, and generate **text sequences**
- Goals:
 - Learn context-dependent representations
 - Capture long-range dependencies
 - Handle complex relationships among large text units
- Sequence modeling architectures are based on deep neural networks (DNNs)!
 - Language exhibits hierarchical structures (e.g., letters form words, words form phrases, phrases form sentences)
 - DNNs learn multiple levels of abstraction across layers, allowing them to capture low-level patterns (e.g., word relations) in lower layers and high-level patterns (e.g., sentence meanings) in higher layers
 - Each layer in DNNs refines the word representations by considering contexts at different granularities (shorter & longer-range contexts), allowing for contextualized understanding of words and sequences

(Recap) Sequence Modeling Architectures



RNN neural networks:

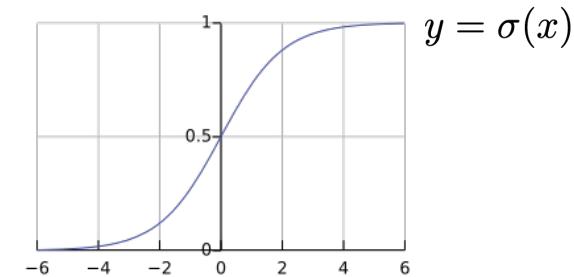
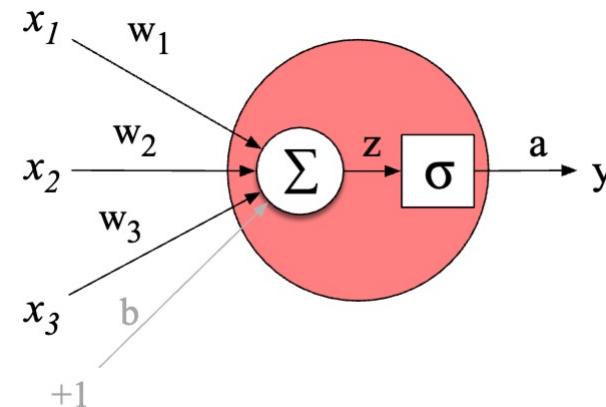
<https://web.stanford.edu/class/cs224n/slides/cs224n-spr2024-lecture05-rnnlm.pdf>



Transformer: <https://arxiv.org/pdf/1706.03762>

(Recap) Neural Network: Basic Unit (Perceptron)

- Input: $\mathbf{x} = [x_1, x_2, x_3]$
- Model parameters (weights & bias): $\mathbf{w} = [w_1, w_2, w_3]$ & b
- Linear computation: $z = \mathbf{w} \cdot \mathbf{x} + b$
- Nonlinear activation: $a = \sigma(z)$



(Recap) Basic Unit (Perceptron): Example

- Input: $\mathbf{x} = [0.5, 0.6, 0.1]$
- Model parameters (weights & bias): $\mathbf{w} = [0.2, 0.3, 0.9]$ & $b = 0.5$
- Linear computation: $z = \mathbf{w} \cdot \mathbf{x} + b = 0.87$
- Nonlinear activation: $a = \sigma(z) = \frac{1}{1 + \exp(-0.87)} \approx 0.70$

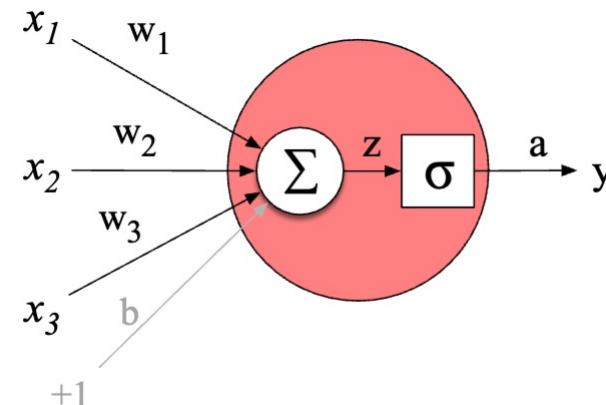
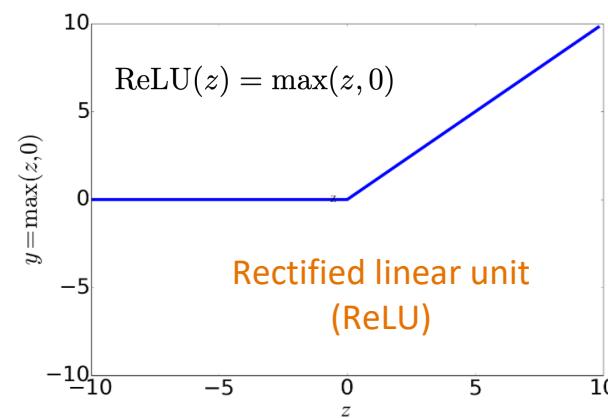
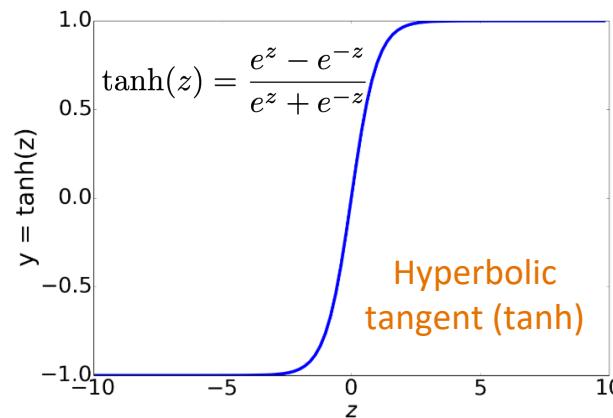


Figure source: <https://web.stanford.edu/~jurafsky/slp3/6.pdf>

(Recap) Common Non-linear Activations

- Why non-linear activations?
- Stacking linear operations will only result in another linear operation
- We wish our network to model complex, non-linear relationships between inputs and outputs



(Recap) Feedforward Network (FFN)

- Feedforward network (FFN) = multi-layer network where the outputs from units in each layer are passed to units in the next higher layer
- FFNs are also called multi-layer perceptrons (MLPs)
- Model parameters in each layer in FFNs: a weight matrix \mathbf{W} and a bias vector \mathbf{b}
 - Each layer has multiple hidden units
 - Recall: a single hidden unit has a weight vector and a bias parameter
 - Weight matrix: combining the weight vector for each unit
 - Bias vector: combining the bias for each unit

(Recap) Example: 2-layer FFN

- Input: $\mathbf{x} = [x_1, x_2, \dots, x_{n_0}]$
- Model parameters (weights & bias): $\mathbf{W} \in \mathbb{R}^{n_1 \times n_0}$, $\mathbf{U} \in \mathbb{R}^{n_2 \times n_1}$ & $\mathbf{b} \in \mathbb{R}^{n_1}$
- Forward computation:

First layer: $\mathbf{h} = \sigma(\mathbf{W}\mathbf{x} + \mathbf{b})$



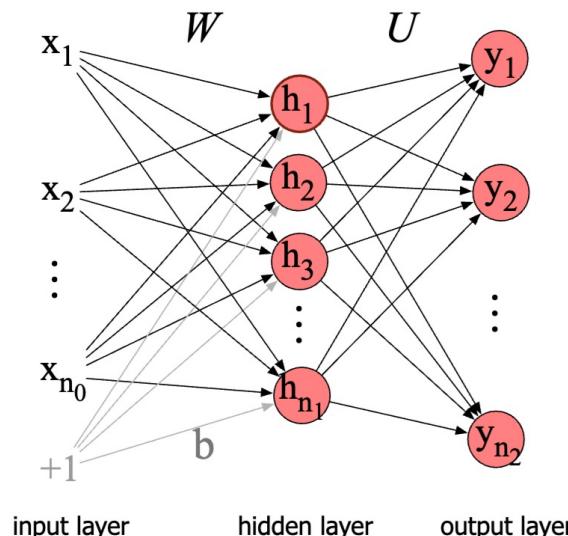
Non-linear function (element-wise)

Second layer: $\mathbf{z} = \mathbf{U}\mathbf{h}$

Output: $\mathbf{y} = \text{softmax}(\mathbf{z})$

Convert to probability distribution

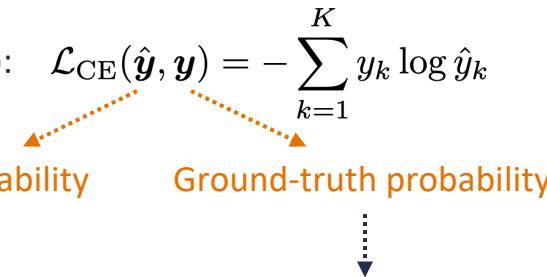
$$= \left[\frac{\exp(z_1)}{\sum_{j=1}^{n_2} \exp(z_j)}, \dots, \frac{\exp(z_{n_2})}{\sum_{j=1}^{n_2} \exp(z_j)} \right]$$



(Recap) Training Objective

- We'll need a **loss function** that models the distance between the model output and the gold/desired output
- The common loss function for classification tasks is **cross-entropy** (CE) loss

K-way classification (K classes): $\mathcal{L}_{\text{CE}}(\hat{\mathbf{y}}, \mathbf{y}) = - \sum_{k=1}^K y_k \log \hat{y}_k$



Model output probability Ground-truth probability

Usually a one-hot vector (one dimension is 1; others are 0): $\mathbf{y} = [0, \dots, 1, \dots, 0]$

$$\mathcal{L}_{\text{CE}}(\hat{\mathbf{y}}, \mathbf{y}) = -\log \hat{y}_c = -\log \frac{\exp(z_c)}{\sum_{j=1}^K \exp(z_j)}$$

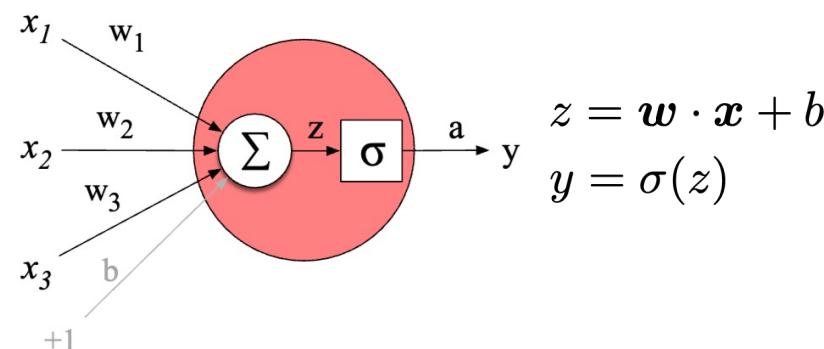
Also called “negative log likelihood (NLL) loss”

↓

c is the ground-truth class

(Recap) Model Training (Forward Pass)

- Most optimization methods for DNNs are based on gradient descent
- First, randomly initialize model parameters
- In each optimization step, run two passes
 - **Forward pass:** evaluate the loss function given the input and current model parameters



(Recap) Model Training (Backward Pass)

- Most optimization methods for DNNs are based on gradient descent
- First, randomly initialize model parameters
- In each optimization step, run two passes
 - **Forward pass:** evaluate the loss function given the input and current model parameters
 - **Backward pass:** update the parameters following the opposite direction of the gradient

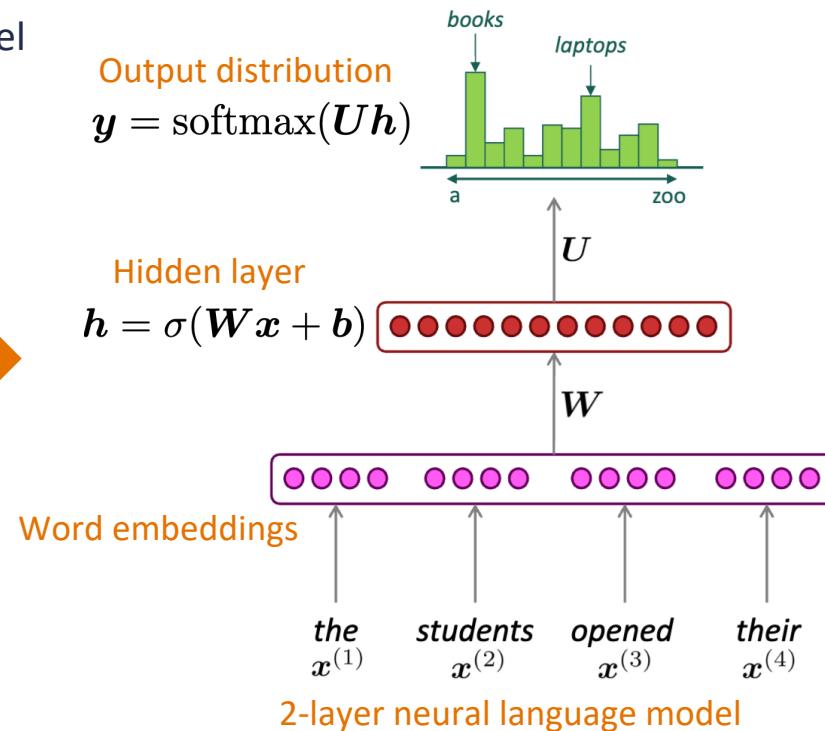
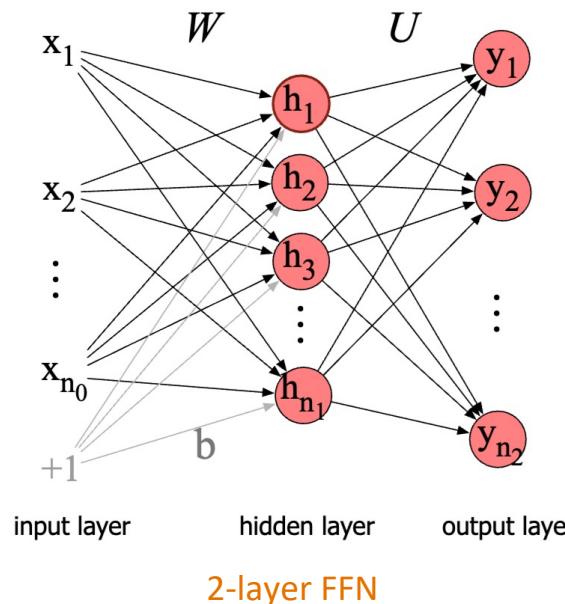
$$\mathbf{w}^{(t+1)} \leftarrow \mathbf{w}^{(t)} - \eta \nabla_{\mathbf{w}} \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y})$$

- Gradient computed via the chain rule $\nabla_{\mathbf{w}} \mathcal{L}(\hat{\mathbf{y}}, \mathbf{y}) = \frac{\partial \mathcal{L}}{\partial \mathbf{w}} = \frac{\partial \mathcal{L}}{\partial \mathbf{y}} \frac{\partial \mathbf{y}}{\partial \mathbf{z}} \frac{\partial \mathbf{z}}{\partial \mathbf{w}}$

Gradient computation taken care of by deep learning libraries
(e.g., PyTorch)

(Recap) Simple Neural Language Model

Instantiate FFN as a neural language model

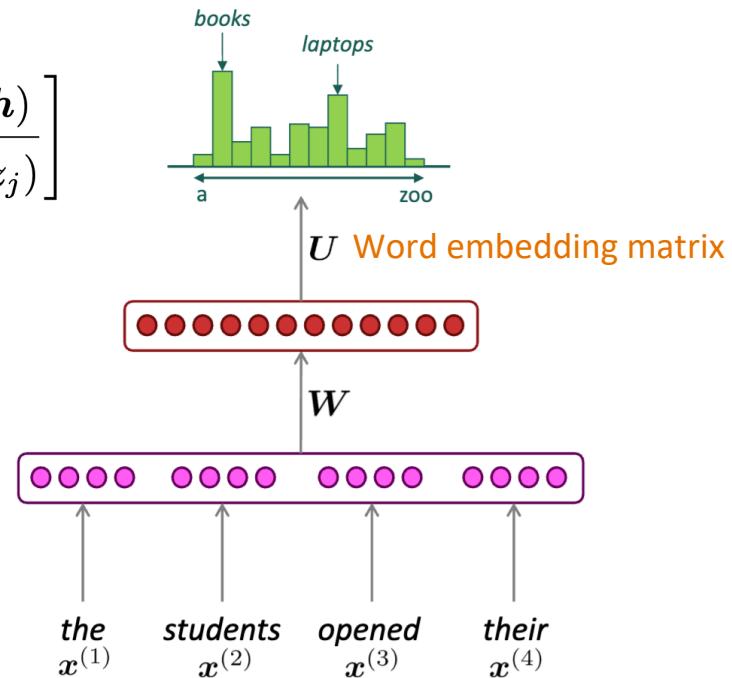


(Recap) Benefits of Neural Language Models

Output distribution

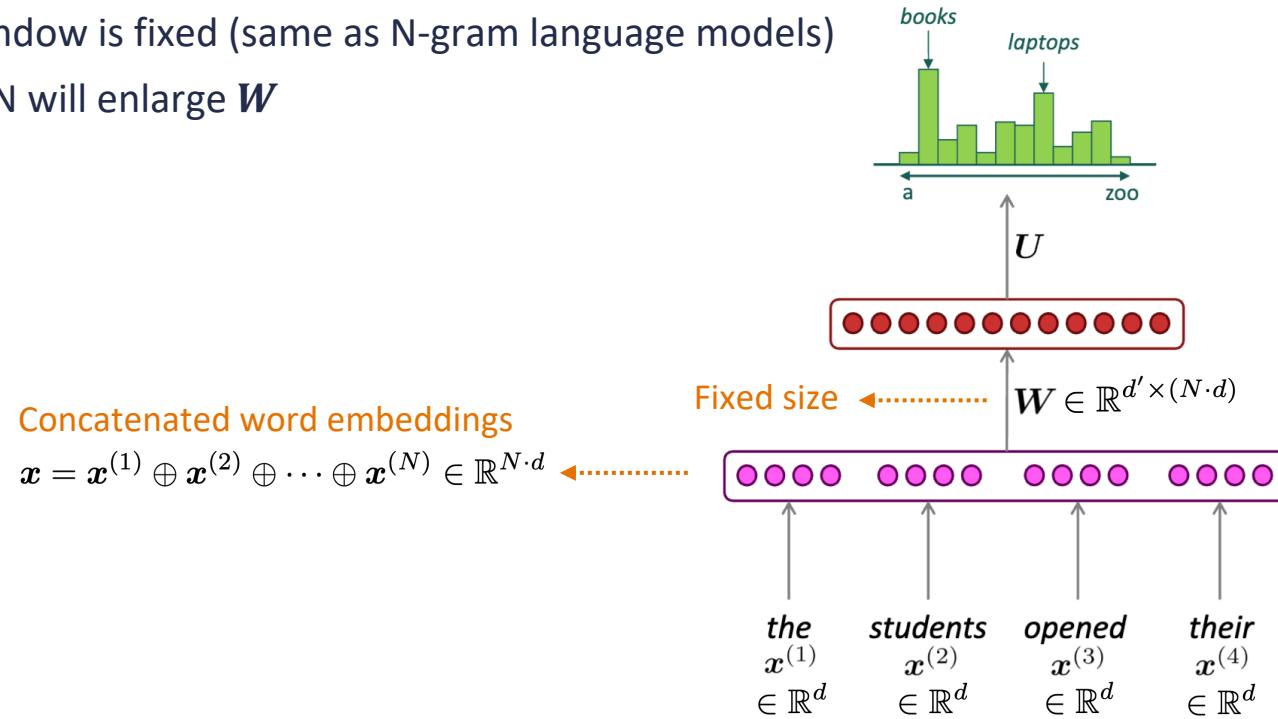
$$\mathbf{y} = \text{softmax}(\mathbf{U}\mathbf{h}) = \left[\underbrace{\frac{\exp(\mathbf{u}_1 \cdot \mathbf{h})}{\sum_{j=1}^{|\mathcal{V}|} \exp(z_j)}, \dots, \frac{\exp(\mathbf{u}_{|\mathcal{V}|} \cdot \mathbf{h})}{\sum_{j=1}^{|\mathcal{V}|} \exp(z_j)}}_{|\mathcal{V}|\text{-dimensions}} \right]$$

- Address sparsity issue:
 - Strictly positive probability on every token in the vocabulary
 - Semantically similar words tend to have similar probabilities



(Recap) Limitations of (Simple) Neural Language Models

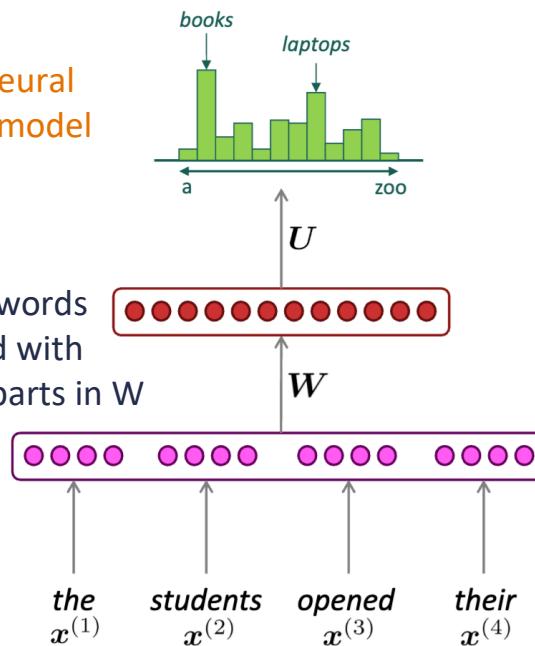
- Context window is fixed (same as N-gram language models)
- Increasing N will enlarge W



(Recap) Recurrent Neural Network (RNN) Overview

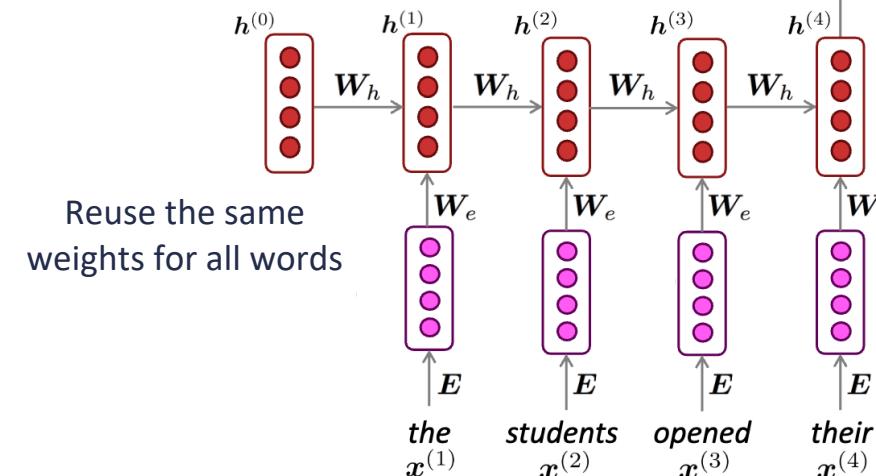
A neural language model that can process inputs of arbitrary lengths

Simple neural language model



Different words multiplied with different subparts in W

Recurrent neural language model

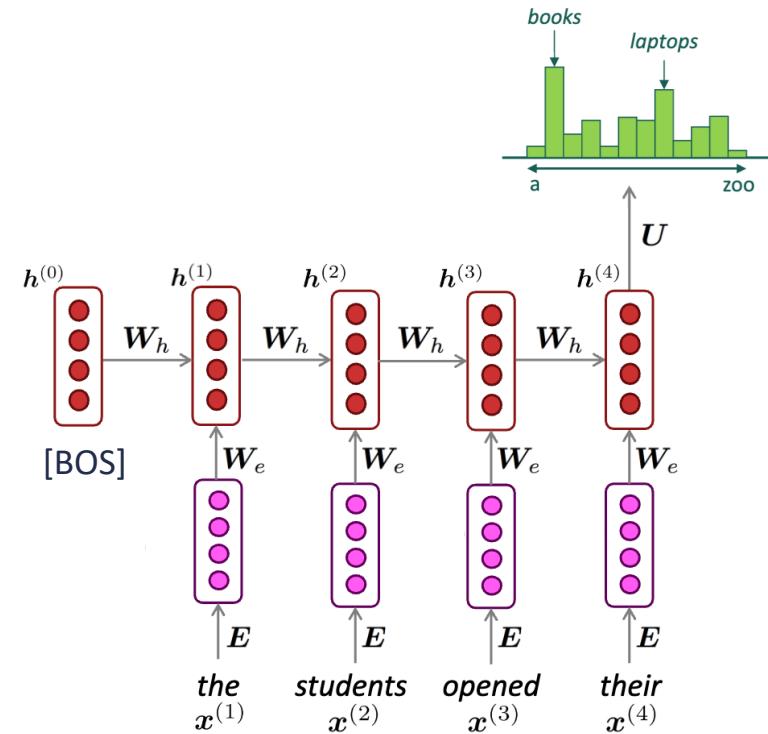


(Recap) RNN Computation

- Hidden states in RNNs are computed based on
 - The hidden state at the previous step (memory)
 - The word embedding at the current step
- Parameters:
 - W_h : weight matrix for the recurrent connection
 - W_e : weight matrix for the input connection

$$h^{(t)} = \sigma \left(W_h h^{(t-1)} + W_e x^{(t)} \right)$$

Hidden states at the
previous word (time step)
Word embedding of the
current word (time step)



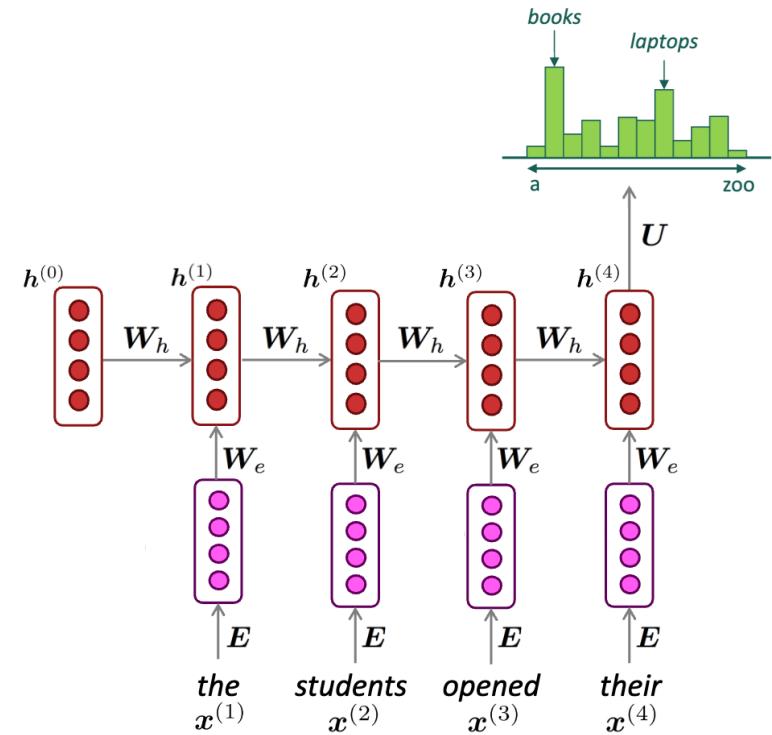
(Recap) RNN Computation

- Input: $\mathbf{x} = [\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}]$
- Initialize $\mathbf{h}^{(0)}$
- For each time step (word) in the input:
 - Compute hidden states:

$$\mathbf{h}^{(t)} = \sigma \left(\mathbf{W}_h \mathbf{h}^{(t-1)} + \mathbf{W}_e \mathbf{x}^{(t)} \right)$$

- Compute output:

$$\mathbf{y}^{(t)} = \text{softmax} \left(\mathbf{U} \mathbf{h}^{(t)} \right)$$



(Recap) RNN Weight Tying

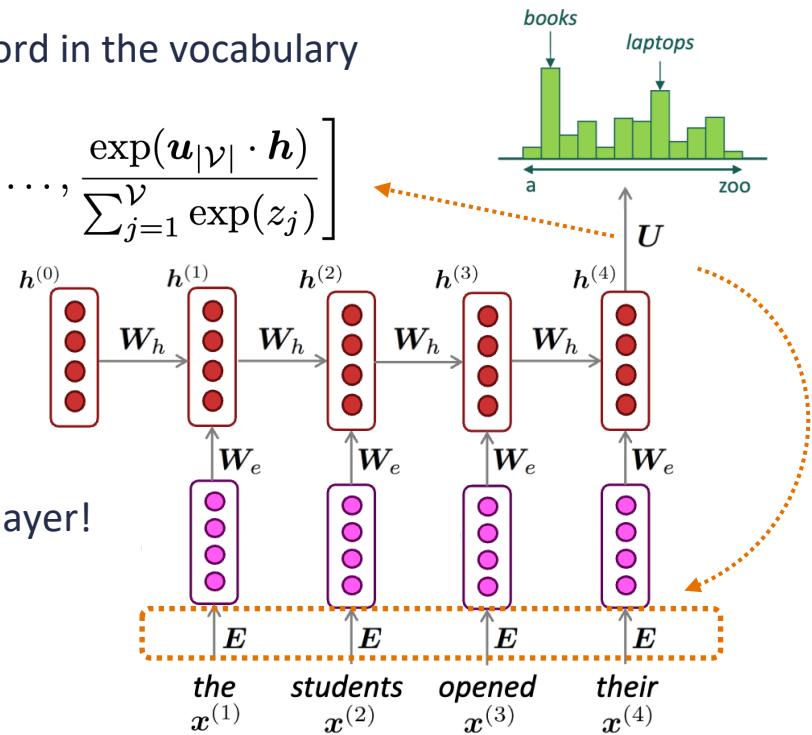
- Role of matrix \mathbf{U} : score the likelihood of each word in the vocabulary

$$\mathbf{y} = \text{softmax}(\mathbf{U}\mathbf{h}) = \left[\frac{\exp(\mathbf{u}_1 \cdot \mathbf{h})}{\sum_{j=1}^{|V|} \exp(z_j)}, \dots, \frac{\exp(\mathbf{u}_{|V|} \cdot \mathbf{h})}{\sum_{j=1}^{|V|} \exp(z_j)} \right]$$

$$\mathbf{U} \in \mathbb{R}^{|V| \times d}$$

Same dimensionality of the word embedding matrix!

- Use the same input embeddings in the softmax layer!
- Weight tying benefits:
 - Improve learning efficiency & effectiveness
 - Reduce the number of parameters in the model



(Recap) RNN for Language Modeling

- Recall that language modeling predicts the next word given previous words

$$p(\mathbf{x}) = p(x^{(1)}) p(x^{(2)}|x^{(1)}) \cdots p(x^{(n)}|x^{(1)}, \dots, x^{(n-1)}) = \prod_{t=1}^n p(x^{(t)}|x^{(1)}, \dots, x^{(t-1)})$$

- How to use RNNs to represent $p(x^{(t)}|x^{(1)}, \dots, x^{(t-1)})$?

Output probability at $(t-1)$ step: $\mathbf{y}^{(t-1)} = \text{softmax}(\mathbf{U}\mathbf{h}^{(t-1)}) := f(\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(t-2)}, \mathbf{x}^{(t-1)})$

$\mathbf{h}^{(t-1)}$ is a function of $\mathbf{h}^{(t-2)}$ and $\mathbf{x}^{(t-1)}$: $\mathbf{h}^{(t-1)} = \sigma(\mathbf{W}_h \mathbf{h}^{(t-2)} + \mathbf{W}_e \mathbf{x}^{(t-1)}) := g(\mathbf{h}^{(t-2)}, \mathbf{x}^{(t-1)})$

$\mathbf{h}^{(t-2)}$ is a function of $\mathbf{h}^{(t-3)}$ and $\mathbf{x}^{(t-2)}$: $\mathbf{h}^{(t-2)} = \sigma(\mathbf{W}_h \mathbf{h}^{(t-3)} + \mathbf{W}_e \mathbf{x}^{(t-2)}) := g(\mathbf{h}^{(t-3)}, \mathbf{x}^{(t-2)})$

⋮

⋮

$\mathbf{h}^{(1)}$ is a function of $\mathbf{h}^{(0)}$ and $\mathbf{x}^{(1)}$: $\mathbf{h}^{(1)} = \sigma(\mathbf{W}_h \mathbf{h}^{(0)} + \mathbf{W}_e \mathbf{x}^{(1)}) := g(\mathbf{h}^{(0)}, \mathbf{x}^{(1)})$

(Recap) RNN Language Model Training

Train the output probability at each time step to predict the next word

$$\mathcal{L}_{\text{LM}}(\mathbf{x}) = \frac{1}{n} \sum_{t=1}^n \mathcal{L}_{\text{CE}}(\hat{\mathbf{y}}^{(t)}, \mathbf{y}^{(t)}) = \frac{1}{n} \sum_{t=1}^n -\log \hat{y}_{x^{(t)}}^{(t)} = \frac{1}{n} \sum_{t=1}^n -\log \frac{\exp(x^{(t)})}{\sum_{w' \in \mathcal{V}} \exp(w')}$$

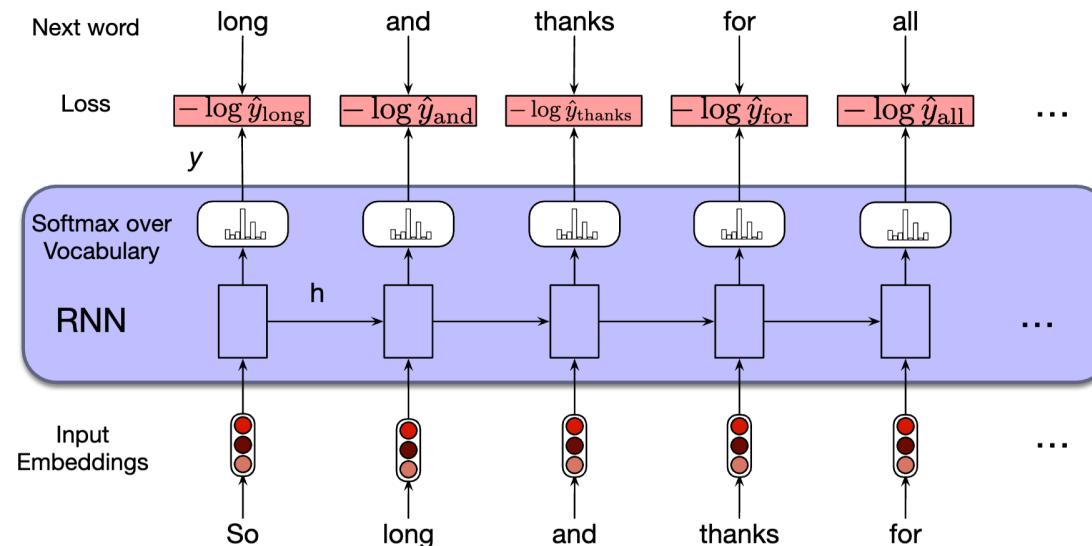


Figure source: <https://web.stanford.edu/~jurafsky/slp3/13.pdf>

(Recap) RNN for Text Generation

- Input [BOS] (beginning-of-sequence) token to the model
- Sample a word from the softmax distribution at the first time step
- Use the word embedding of that first word as the input at the next time step
- Repeat until the [EOS] (end-of-sequence) token is generated

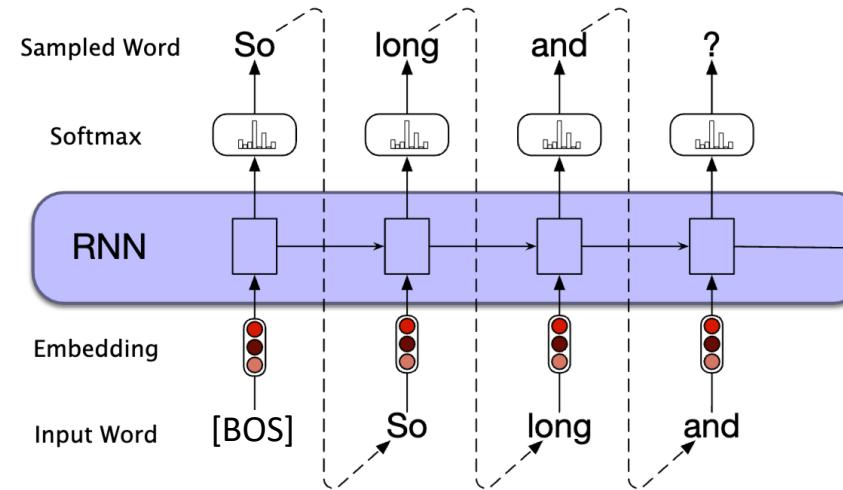


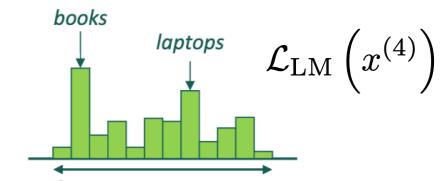
Figure source: <https://web.stanford.edu/~jurafsky/slp3/13.pdf>

Agenda

- RNN Limitations
- Advanced RNNs
- Transformer Overview
- Self-Attention

Vanishing & Exploding Gradient

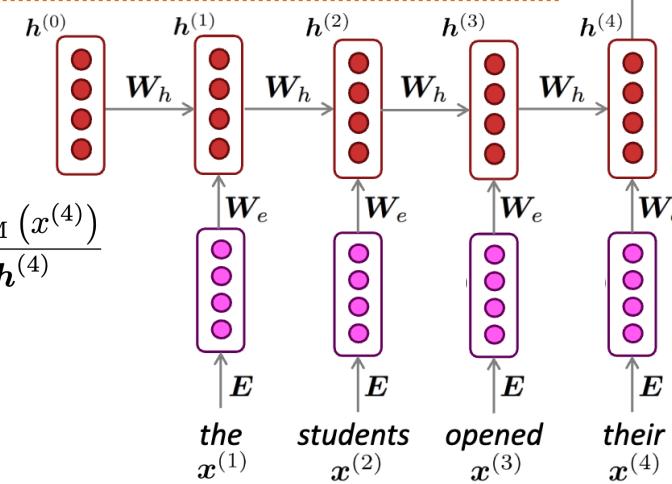
- Gradient signal from far away can be unstable!
- Vanishing gradient = many small gradients multiplied together
- Exploding gradient = many large gradients multiplied together



Gradient backpropagation

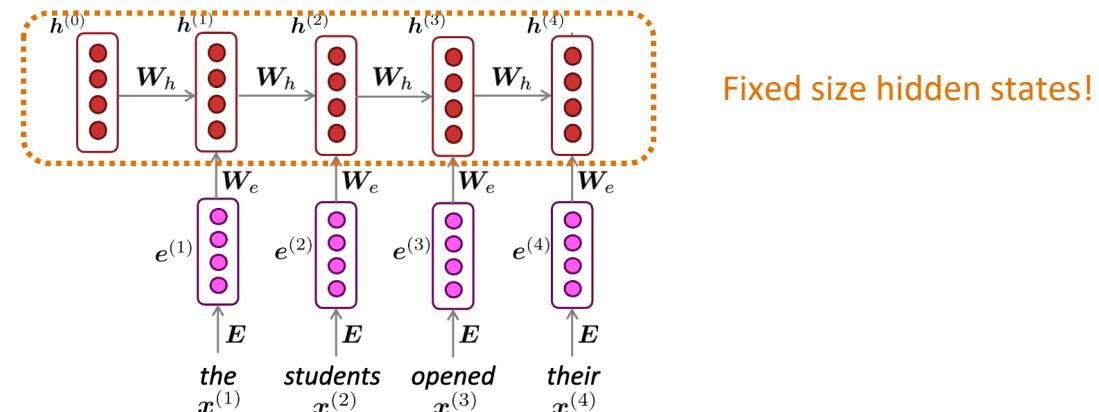
Lots of gradient multiplications!

$$\frac{\partial \mathcal{L}_{\text{LM}}(x^{(4)})}{\partial h^{(0)}} = \boxed{\frac{\partial h^{(1)}}{\partial h^{(0)}} \frac{\partial h^{(2)}}{\partial h^{(1)}} \frac{\partial h^{(3)}}{\partial h^{(2)}} \frac{\partial h^{(4)}}{\partial h^{(3)}}} \frac{\partial \mathcal{L}_{\text{LM}}(x^{(4)})}{\partial h^{(4)}}$$



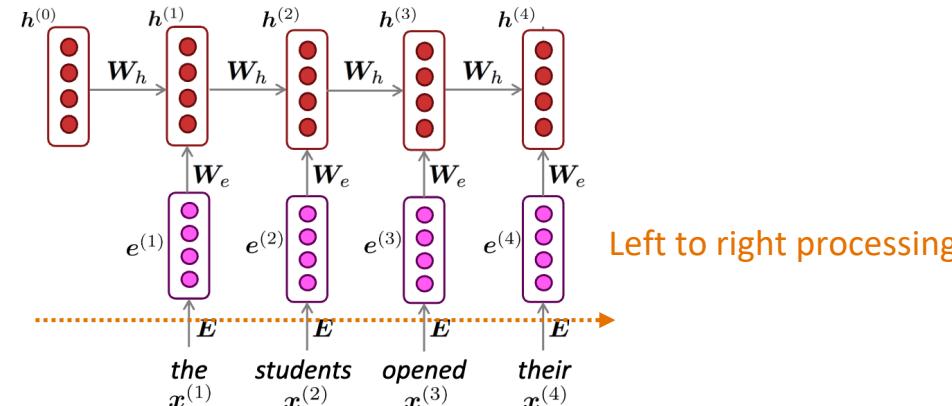
Difficulty in Capturing Long-Term Dependencies

- RNNs are theoretically capable of remembering information over arbitrary lengths of input, but they struggle in practice with long-term dependencies
- RNNs use a fixed-size hidden state to encode an entire sequence of variable length; the hidden state is required to compress a lot of information
- RNNs might give more weight to the most recent inputs and may ignore or “forget” important information at the beginning of the sentence while processing the end



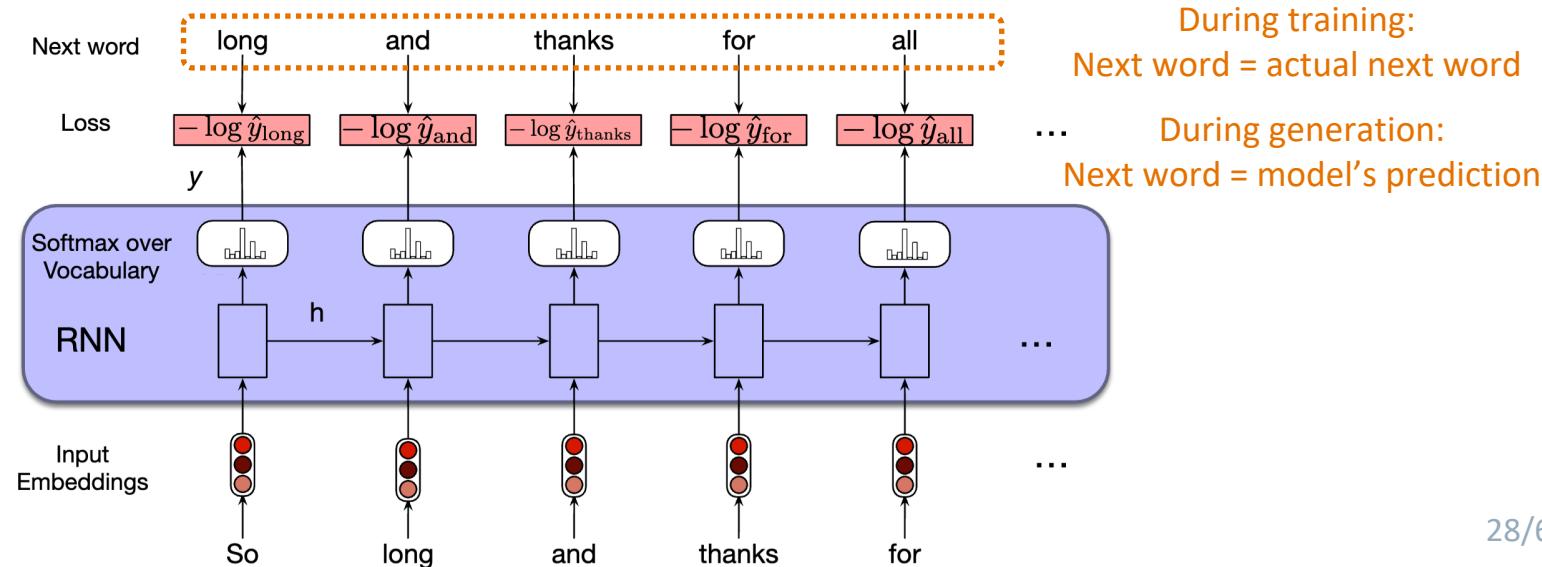
Lack of Bidirectionality

- RNNs process the input sequence step by step from the beginning to the end (left to right for English)
- At each time step, the hidden state only has access to the information from the past without being able to leverage future contexts
- Example: “The bank is on the river” -> the word “bank” can be correctly disambiguated only if the model has access to the word “river” later in the sentence



Exposure Bias

- **Teacher forcing/exposure bias:** during RNN training, the model always receives the **correct** next word from the training data as input for the next step
- When the model has to predict sequences on its own, it may perform poorly if it hasn't learned how to correct its own mistakes



Agenda

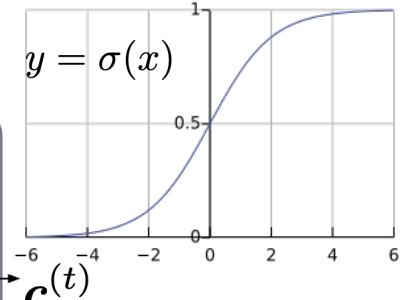
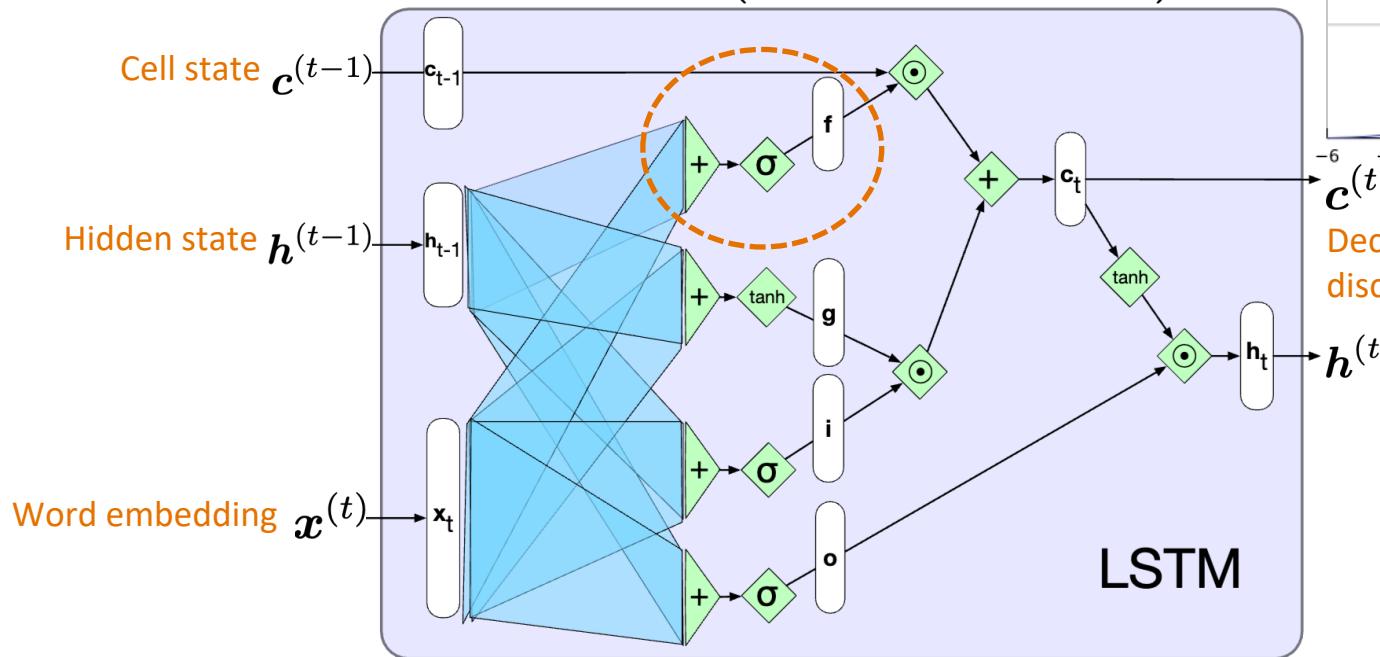
- RNN Limitations
- Advanced RNNs
- Transformer Overview
- Self-Attention

Long Short-Term Memory (LSTM)

- Challenge in RNNs: information encoded in hidden states tends to be local; distant information gets lost
- LSTM design intuition:
 - Remove information no longer needed from the context
 - Add information likely to be needed for future time steps
- Inputs at each time step:
 - Word embedding of the current word
 - Hidden state from the previous time step
 - **Memory/cell state**
- Three gates:
 - Forget gate
 - Add/input gate
 - Output gate

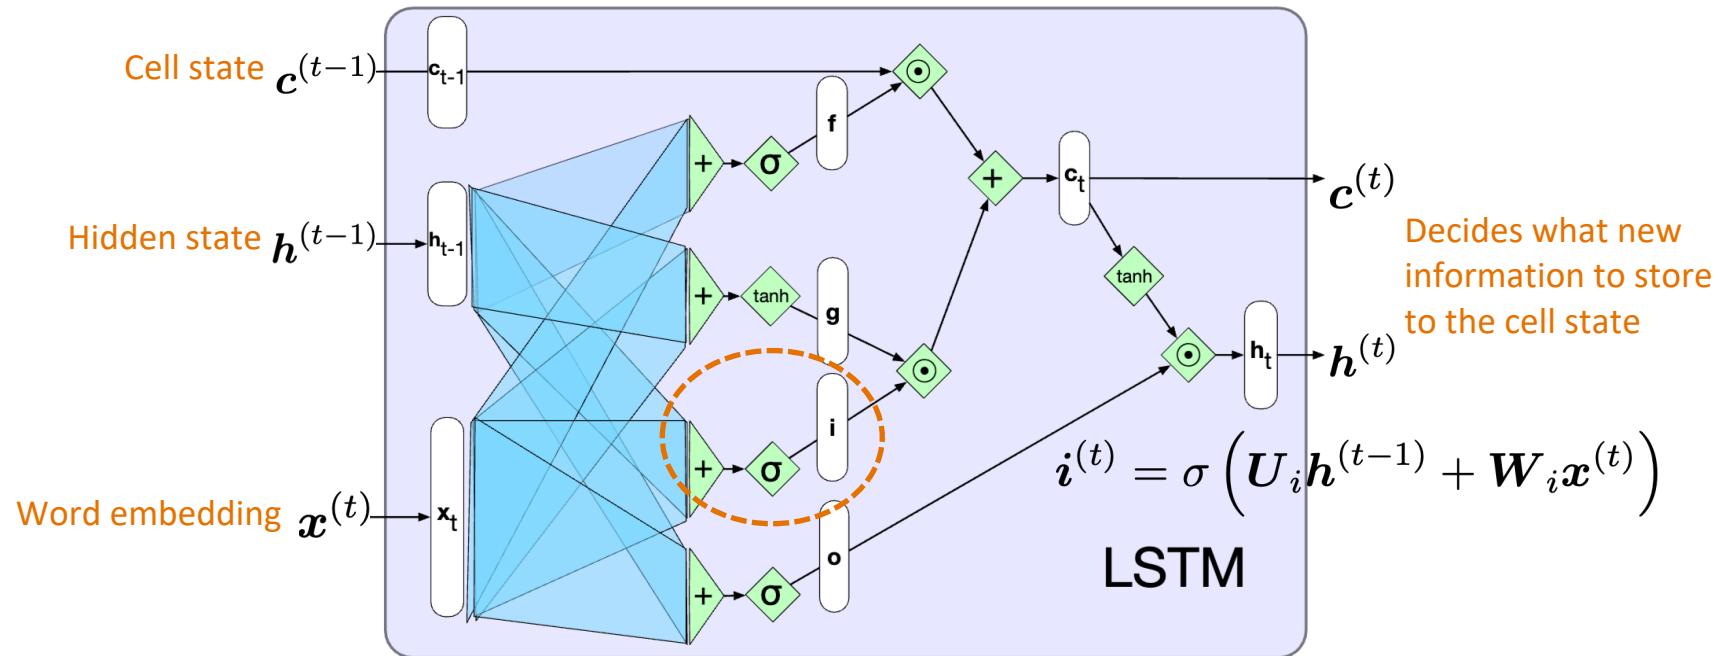
LSTM Computation (Forget Gate)

$$f^{(t)} = \sigma \left(\mathbf{U}_f h^{(t-1)} + \mathbf{W}_f x^{(t)} \right)$$



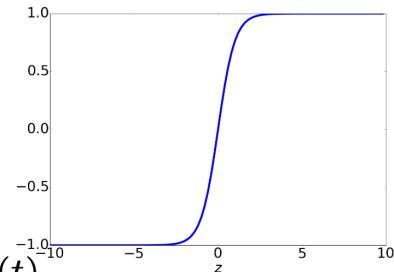
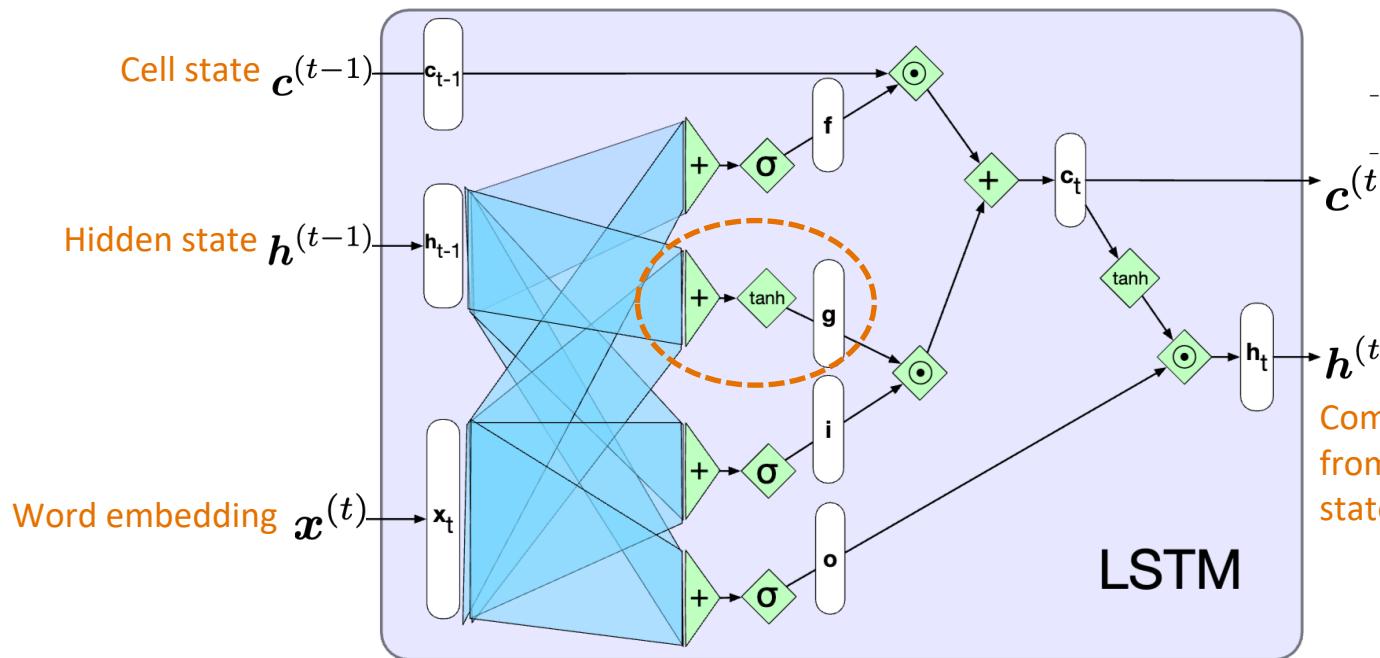
$c^{(t)}$
Decides what information to
discard from the cell state

LSTM Computation (Add/Input Gate)



LSTM Computation (Candidate Cell State)

$$g^{(t)} = \tanh \left(\mathbf{U}_g h^{(t-1)} + \mathbf{W}_g x^{(t)} \right)$$

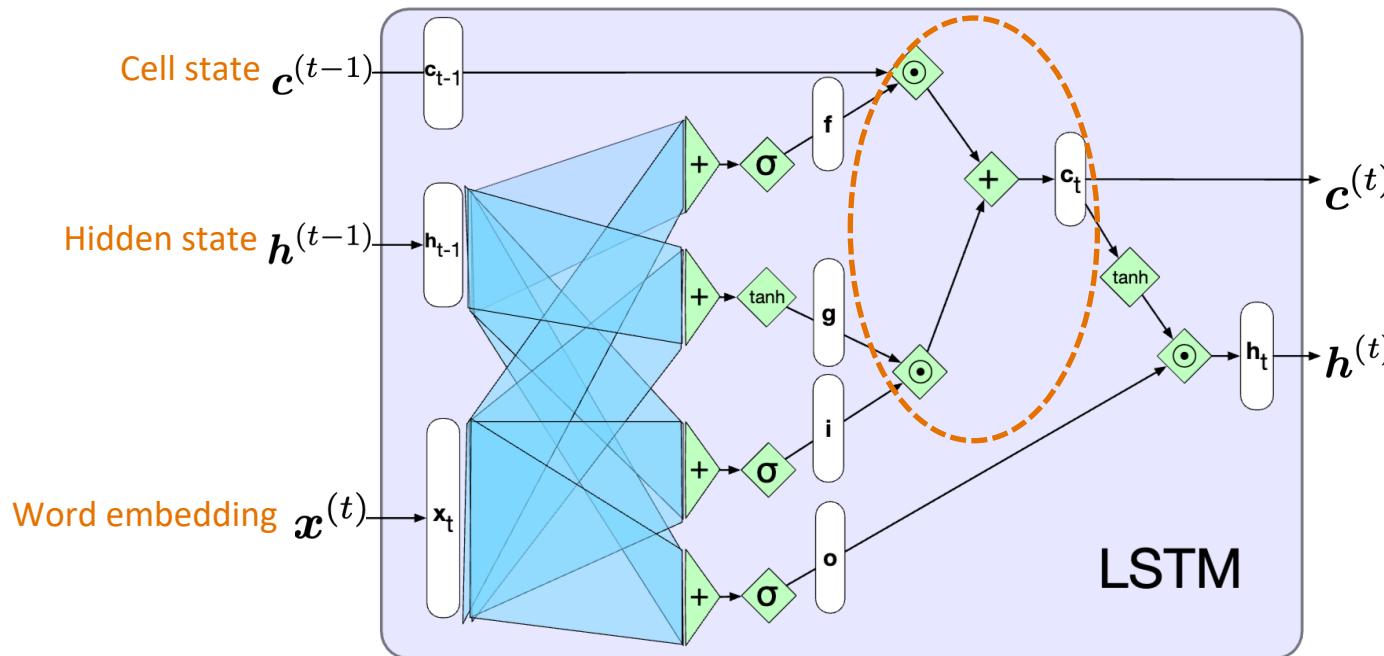


$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

Compute information needed from the previous hidden state and current inputs

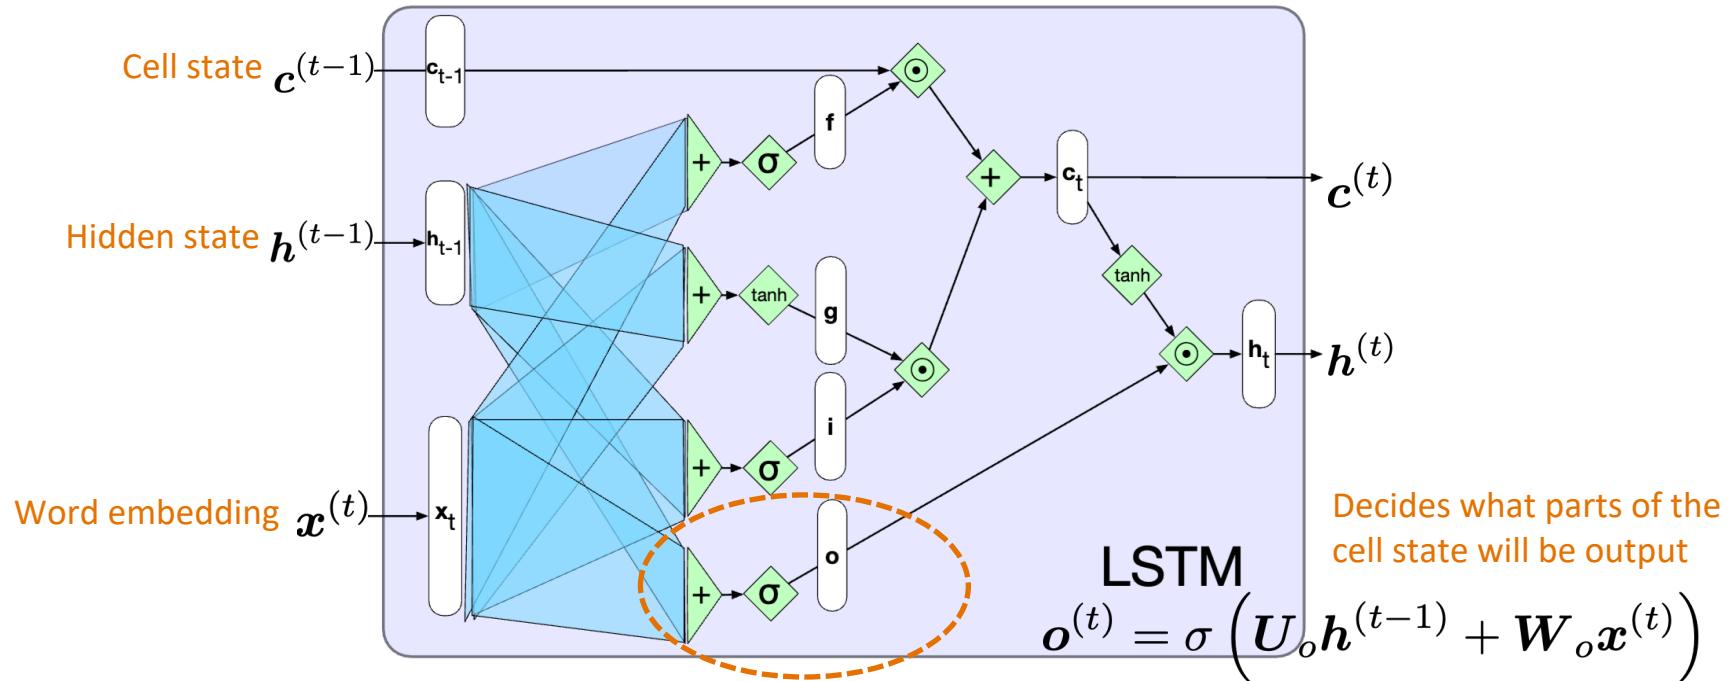
LSTM Computation (Cell State Update)

$$\mathbf{c}^{(t)} = \mathbf{i}^{(t)} \odot \mathbf{g}^{(t)} + \mathbf{f}^{(t)} \odot \mathbf{c}^{(t-1)}$$

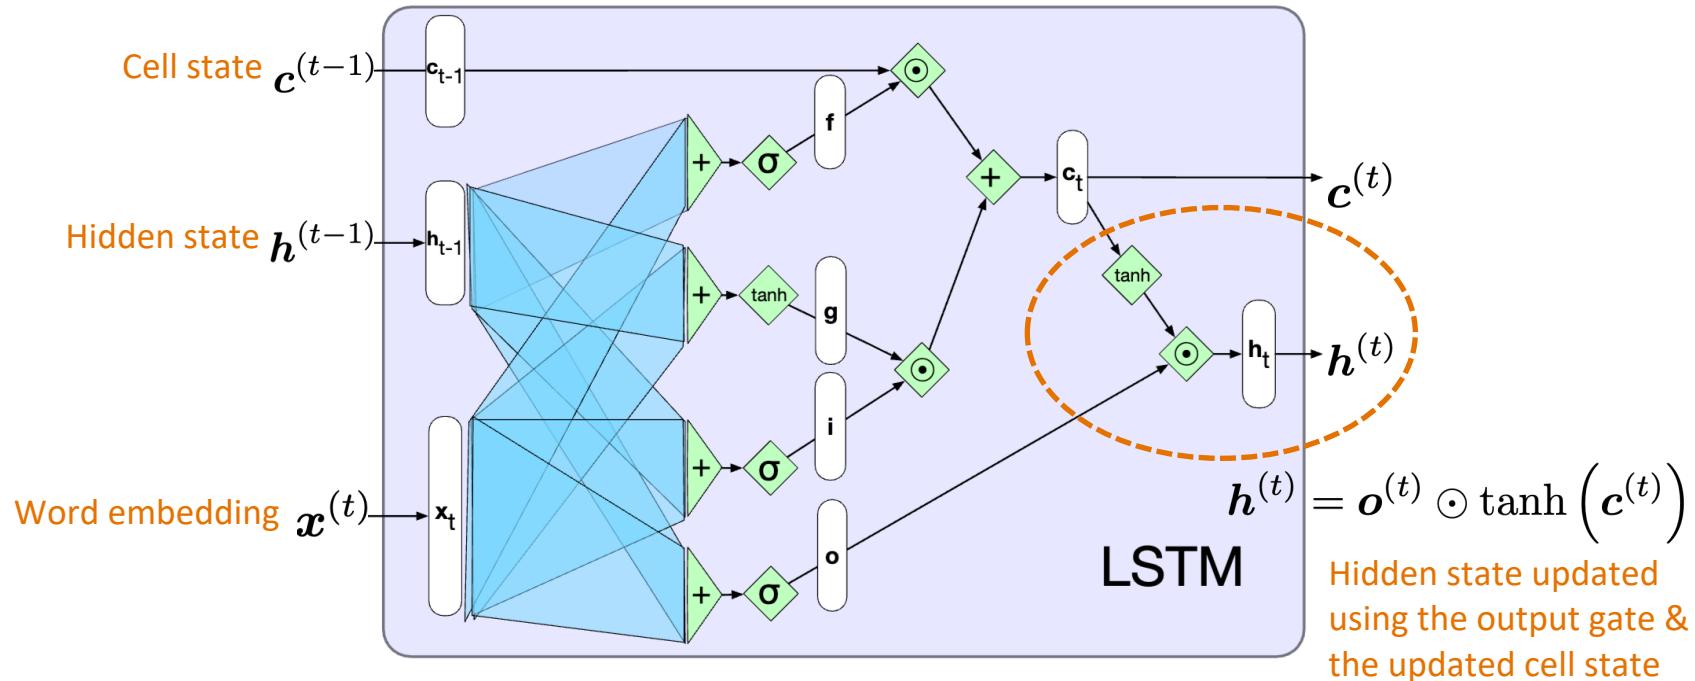


Cell state updated by combining the input gate, candidate cell state, forget gate & previous cell state

LSTM Computation (Output Gate)

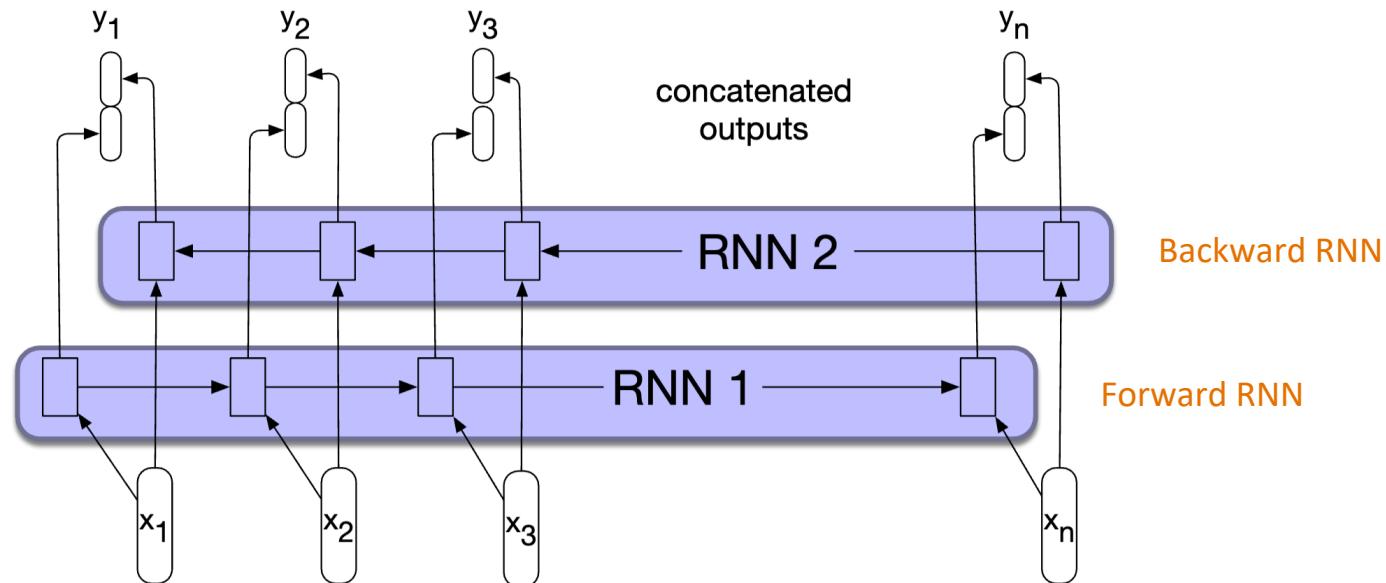


LSTM Computation (Hidden State Update)



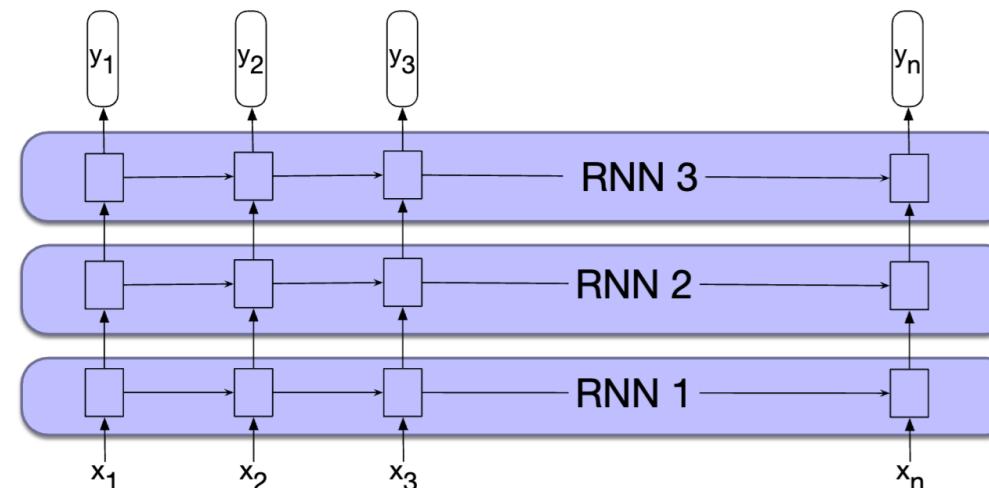
Bidirectional RNNs

- Separate models are trained in the forward and backward directions
- Hidden states from both RNNs are concatenated as the final representations



Deep RNNs

- We can stack multiple RNN layers to build deep RNNs
- The output of a lower level serves as the input to higher levels
- The output of the last layer is used as the final output



Summary: Sequence Modeling

- Sequence modeling goals:
 - Learn context-dependent representations
 - Capture long-range dependencies
 - Handle complex relationships among large text units
- Use deep learning architectures to understand, process, and generate text sequences
- Why DNNs?
 - The multi-layer structure in DNNs mirrors the hierarchical structures in language
 - DNNs learn multiple levels of semantics across layers: low-level patterns (e.g., relations between words) in lower layers & high-level patterns (e.g., sentence meanings) in higher layers

Summary: Neural Language Models

- Address the sparsity issue in N-gram language models by computing the output distribution based on distributed representations (with semantic information)
- Simple neural language models based on feedforward networks suffer from the fixed context window issue
 - Can only model a fixed number of words (similar to N-gram assumption)
 - Increasing the context window requires adding more model parameters

Summary: Recurrent Neural Networks

- General idea: Use the same set of model weights to process all input words
- RNNs as language models
 - Theoretically able to process infinitely long sequences
 - Practically can only keep track of recent contexts
- Training issues: vanishing & exploding gradients
- LSTM is a prominent RNN variant to keep track of both long-term and short-term memories via multiple gates

Agenda

- RNN Limitations
- Advanced RNNs
- Transformer Overview
- Self-Attention

Transformer: Overview

- Transformer is a specific kind of sequence modeling architecture (based on DNNs)
 - Use attention to replace recurrent operations in RNNs
 - The most important architecture for language modeling (almost all LLMs are based on Transformers)!
-

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

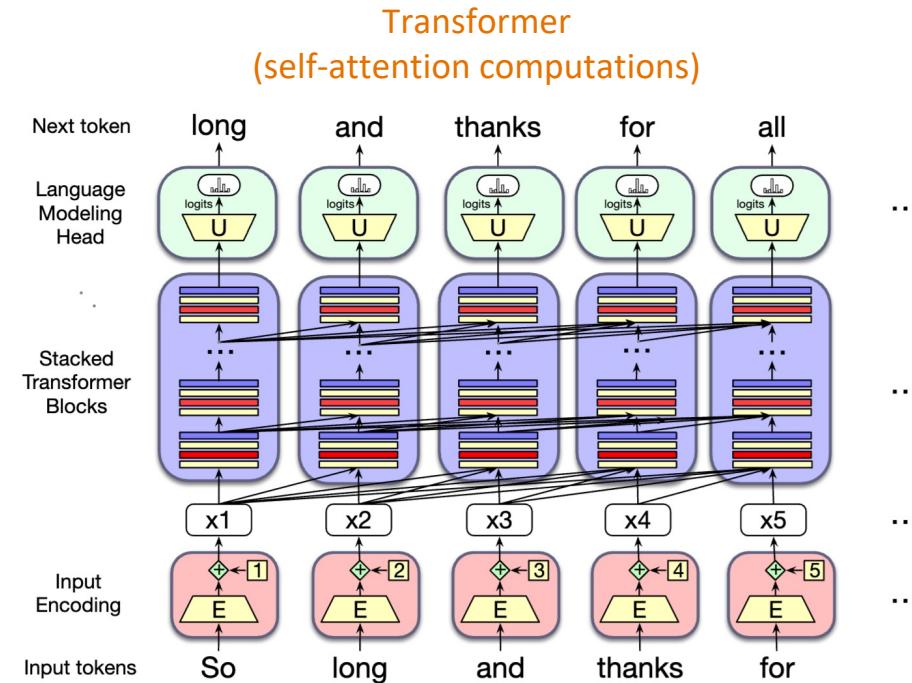
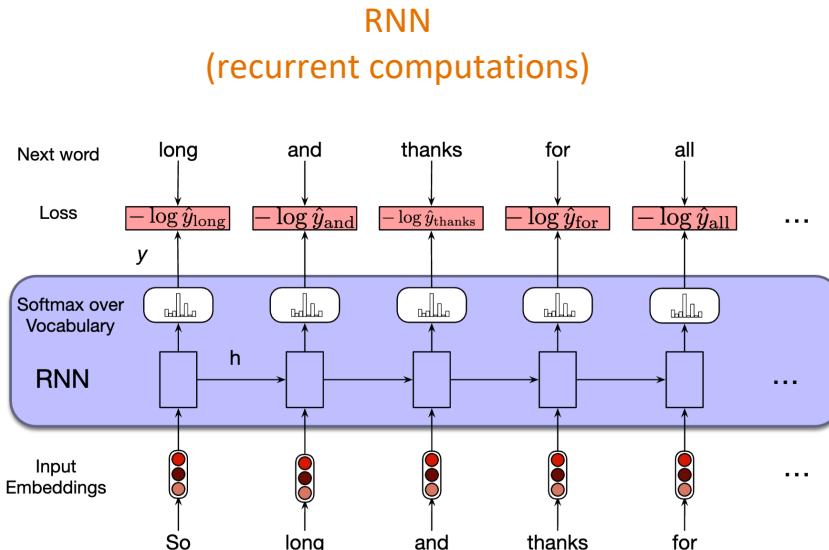
Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Lukasz Kaiser*
Google Brain
lukasz.kaiser@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com

Transformer vs. RNN



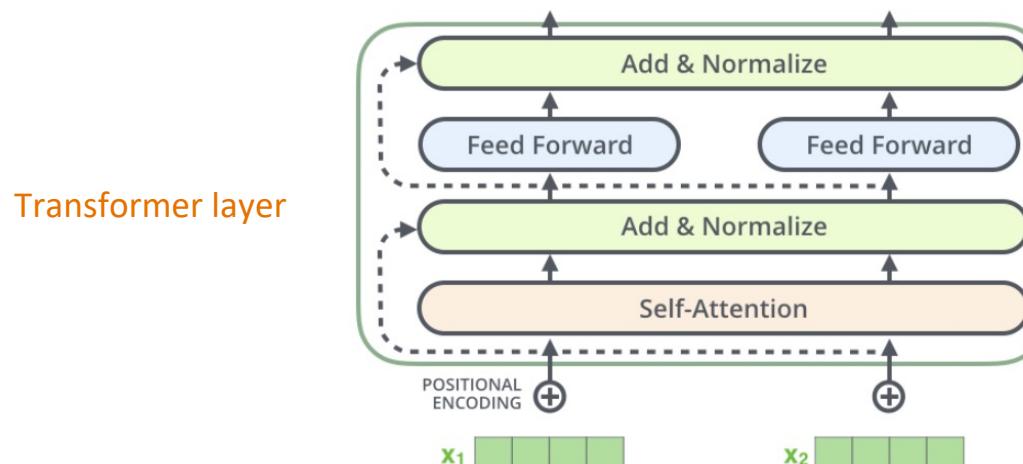
Transformer: Motivation

- Parallel token processing
 - RNN: process one token at a time (computation for each token depends on previous ones)
 - Transformer: process all tokens in a sequence in parallel
- Long-term dependencies
 - RNN: bad at capturing distant relating tokens (vanishing gradients)
 - Transformer: directly access any token in the sequence, regardless of its position
- Bidirectionality
 - RNN: can only model sequences in one direction
 - Transformer: inherently allow bidirectional sequence modeling via attention

Transformer Layer

Each Transformer layer contains the following important components:

- Self-attention
- Feedforward network
- Residual connections + layer norm



Agenda

- RNN Limitations
- Advanced RNNs
- Transformer Overview
- Self-Attention

Self-Attention: Intuition

- Attention: weigh the importance of different words in a sequence when processing a specific word
 - “When I’m looking at this word, which other words should I pay attention to in order to understand it better?”
- **Self-attention:** each word attends to other words in the **same** sequence
- Example: “The quick brown fox jumps over the lazy dog.”
 - Suppose we are learning attention for the word “**jumps**”
 - With self-attention, “**jumps**” can decide which other words in the sentence it should focus on to better understand its meaning
 - Might assign high attention to “fox” (the subject) & “over” (the preposition)
 - Might assign less attention to words like “the” or “lazy”

Self-Attention: Example

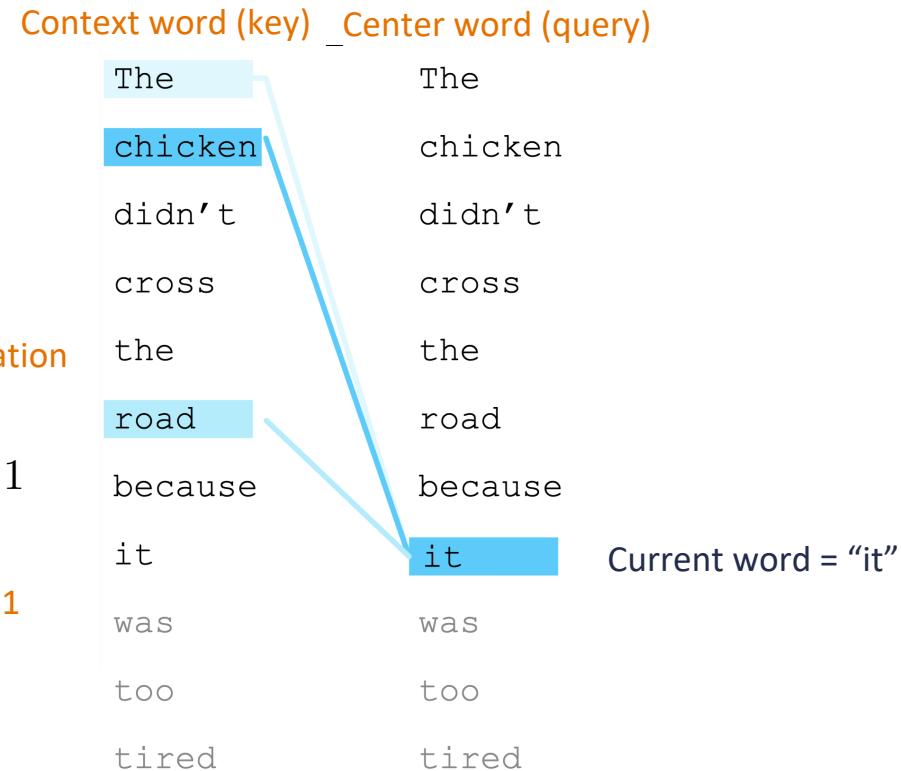
Derive the center word representation as a weighted sum of context representations!

Center word representation

Context word representation

$$\mathbf{a}_i = \sum_{x_j \in \mathbf{x}} \alpha_{ij} \mathbf{x}_j, \quad \sum_{x_j \in \mathbf{x}} \alpha_{ij} = 1$$

Attention score $i \rightarrow j$, summed to 1



Self-Attention: Attention Score Computation

- Attention score is given by the softmax function over vector dot product

$$\mathbf{a}_i = \sum_{x_j \in \mathbf{x}} \alpha_{ij} \mathbf{x}_j, \quad \sum_{x_j \in \mathbf{x}} \alpha_{ij} = 1$$

$$\alpha_{ij} = \text{Softmax}(\mathbf{x}_i \cdot \mathbf{x}_j)$$

Center word (query) representation Context word (key) representation

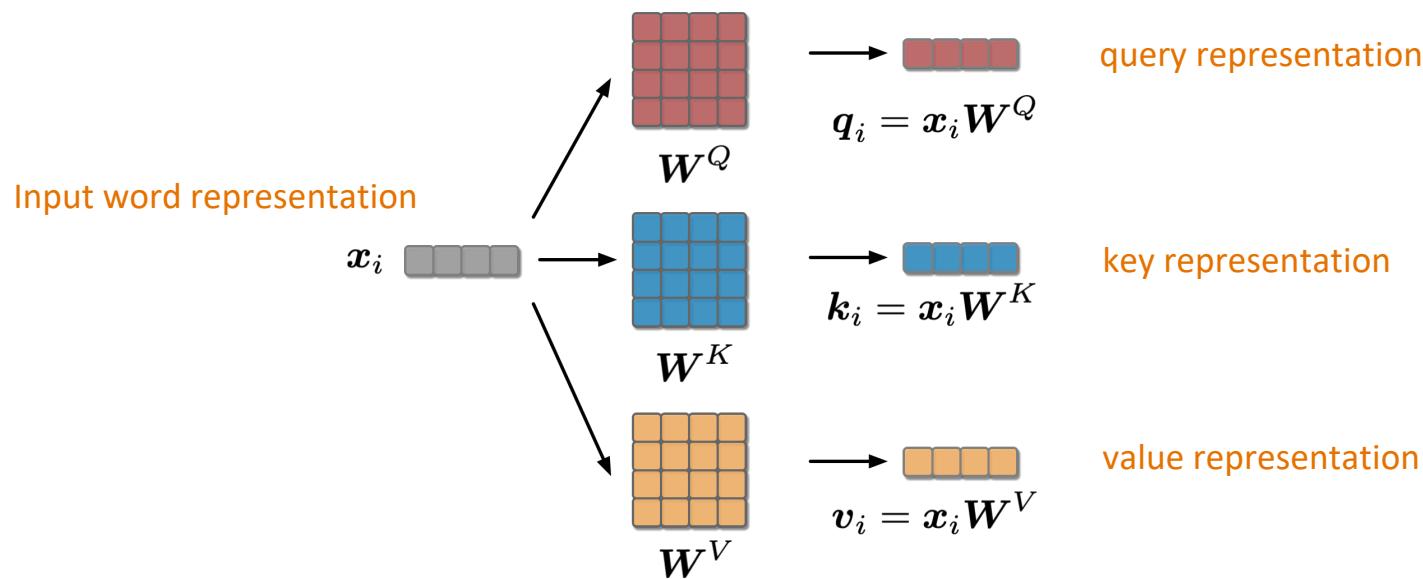
- Why use two copies of word representations for attention computation?
 - We want to reflect the different roles a word plays (as the target word being compared to others, or as the context word being compared to the target word)
 - If using the same copy of representations for attention calculation, a word will (almost) always attend to itself heavily due to high dot product with itself!

Self-Attention: Query, Key, and Value

- Each word in self-attention is represented by three different vectors
 - Allow the model to flexibly capture different types of relationships between tokens
- **Query (Q):**
 - Represent the current word seeking information about
- **Key (K):**
 - Represent the reference (context) against which the query is compared
- **Value (V):**
 - Represent the actual content associated with each token to be aggregated as final output

Self-Attention: Query, Key, and Value

Each self-attention module has three weight matrices applied to the input word vector to obtain the three copies of representations



Self-Attention: Overall Computation

- Input: single word vector of each word \mathbf{x}_i
- Compute Q, K, V representations for each word:

$$\mathbf{q}_i = \mathbf{x}_i \mathbf{W}^Q \quad \mathbf{k}_i = \mathbf{x}_i \mathbf{W}^K \quad \mathbf{v}_i = \mathbf{x}_i \mathbf{W}^V$$

- Compute attention scores with Q and K
 - The dot product of two vectors usually has an expected magnitude proportional to \sqrt{d}
 - Divide the attention score by \sqrt{d} to avoid extremely large values in softmax function

$$\alpha_{ij} = \text{Softmax} \left(\frac{\mathbf{q}_i \cdot \mathbf{k}_j}{\sqrt{d}} \right)$$

.....

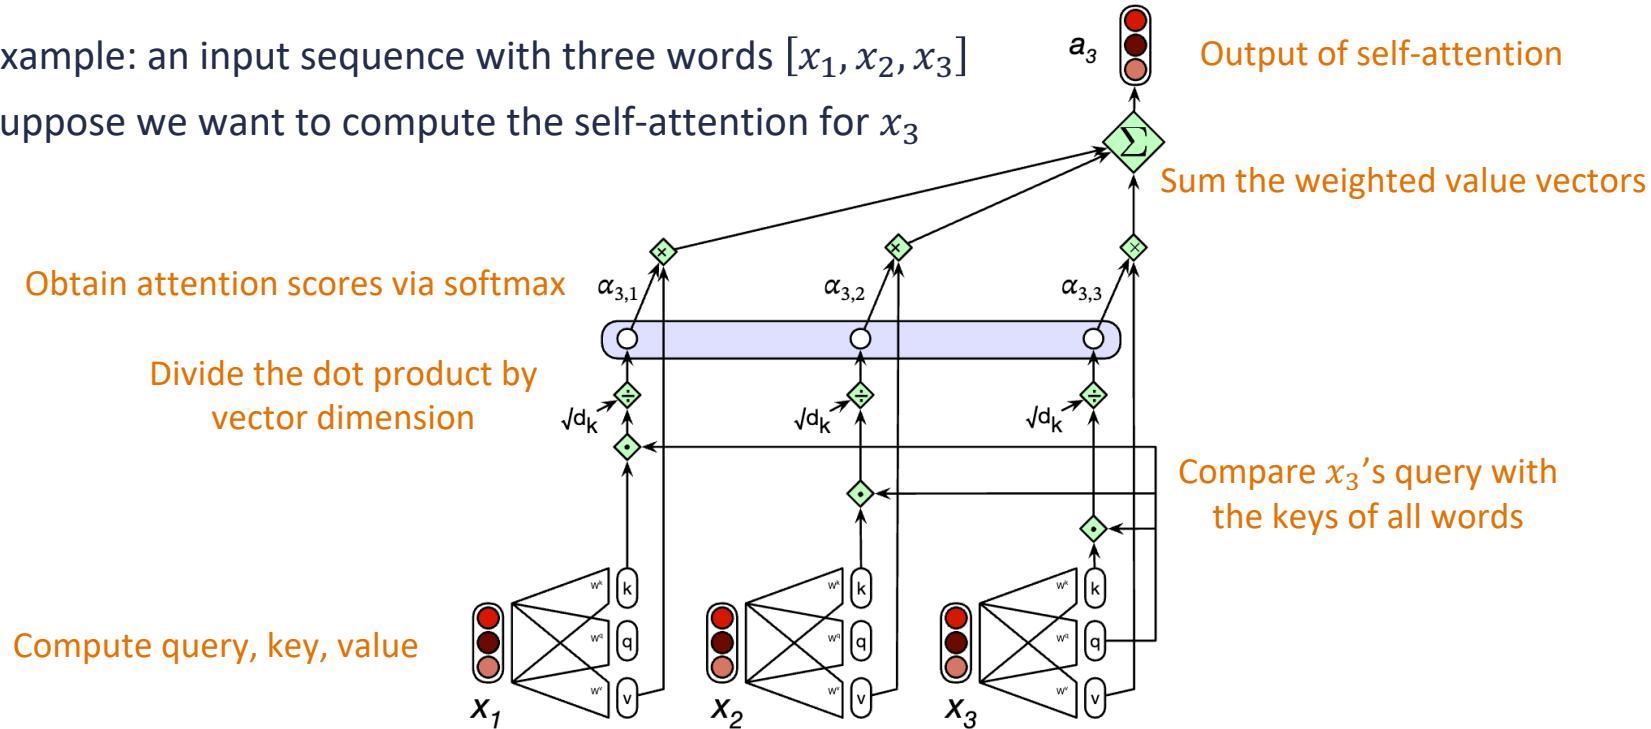
Dimensionality of q and k

- Sum the value vectors weighted by attention scores

$$\mathbf{a}_i = \sum_{x_j \in \mathbf{x}} \alpha_{ij} \mathbf{v}_j$$

Self-Attention: Illustration

- Example: an input sequence with three words $[x_1, x_2, x_3]$
- Suppose we want to compute the self-attention for x_3



Multi-Head Self-Attention

- Transformers use multiple attention heads for each self-attention module
- Intuition:
 - Each head might attend to the context for different purposes (e.g., particular kinds of patterns in the context)
 - Heads might be specialized to represent different linguistic relationships

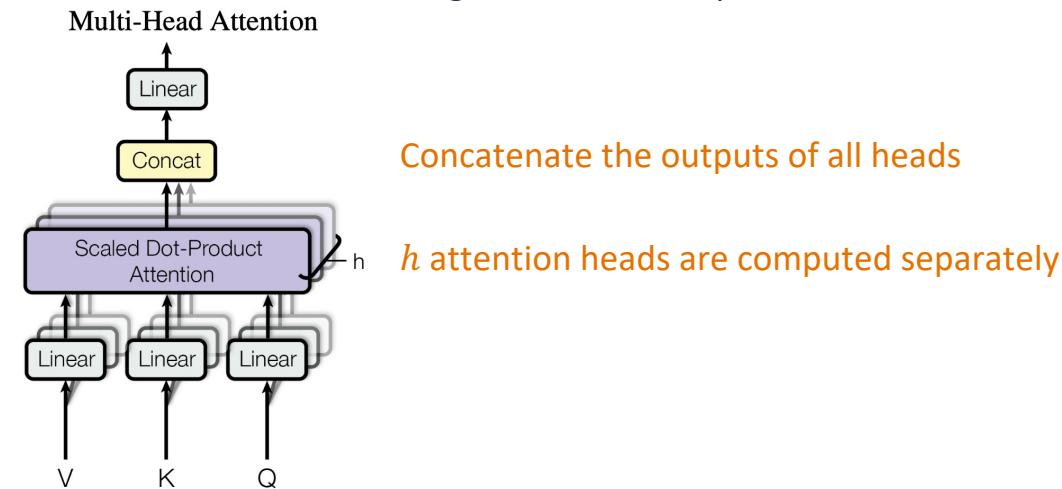


Figure source: <https://arxiv.org/pdf/1706.03762>

Multi-Head Self-Attention Variants

- Multi-query attention ([Fast Transformer Decoding: One Write-Head is All You Need](#)): share keys and values across all attention heads
- Grouped-query attention ([GQA: Training Generalized Multi-Query Transformer Models from Multi-Head Checkpoints](#)): share keys and values within groups of heads

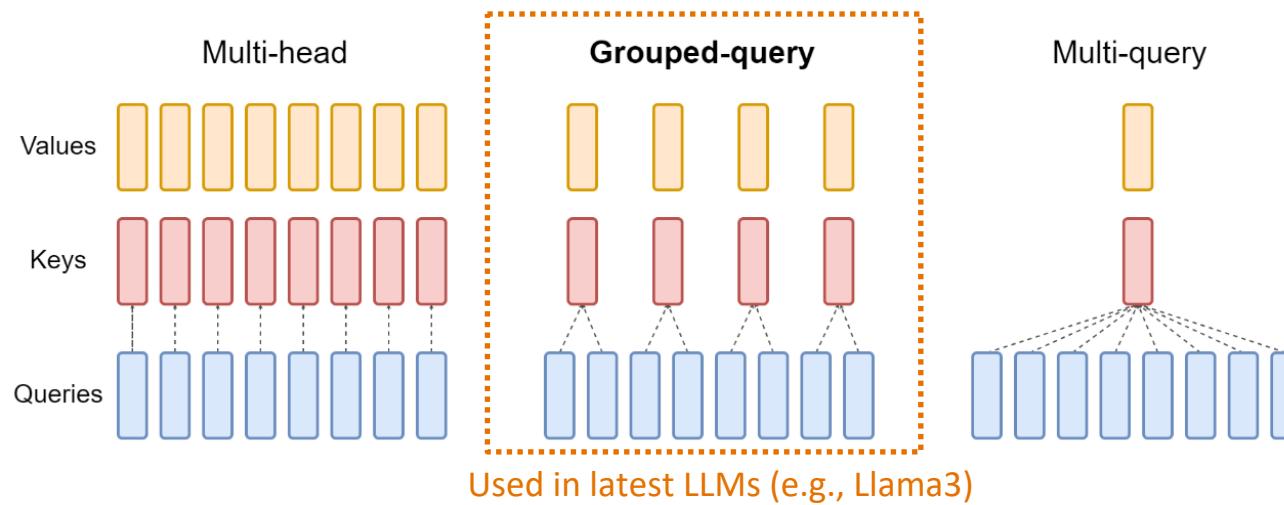


Figure source: <https://arxiv.org/pdf/2305.13245>

Parallel Computation of QKV

- Self-attention computation performed for each token is independent of other tokens
- Easily parallelize the entire computation, taking advantage of the efficient matrix multiplication capability of GPUs
- Process an input sequence with N words in parallel

Compute QKV for one word: $\mathbf{q}_i = \mathbf{x}_i \mathbf{W}^Q \quad \mathbf{k}_i = \mathbf{x}_i \mathbf{W}^K \quad \mathbf{v}_i = \mathbf{x}_i \mathbf{W}^V \in \mathbb{R}^d$

Stacking N input vectors: $\mathbf{Q} = \mathbf{X} \mathbf{W}^Q \quad \mathbf{K} = \mathbf{X} \mathbf{W}^K \quad \mathbf{V} = \mathbf{X} \mathbf{W}^V \in \mathbb{R}^{N \times d}$

$$\mathbf{X} = \begin{bmatrix} \text{---} & \mathbf{x}_1 & \text{---} \\ \text{---} & \mathbf{x}_2 & \text{---} \\ \dots & \dots & \dots \\ \text{---} & \mathbf{x}_N & \text{---} \end{bmatrix}$$

Parallel Computation of Attention

Attention computation can also be written in matrix form

Compute attention for one word: $a_i = \text{Softmax} \left(\frac{\mathbf{q}_i \cdot \mathbf{k}_j}{\sqrt{d}} \right) \cdot \mathbf{v}_j$

Compute attention for one N words: $\mathbf{A} = \text{Softmax} \left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d}} \right) \mathbf{V}$

Attention matrix

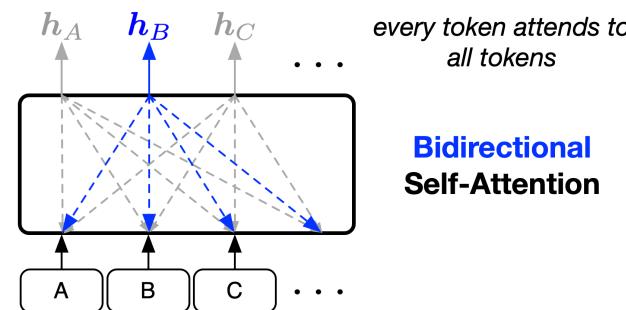
$\mathbf{q1} \cdot \mathbf{k1}$	$\mathbf{q1} \cdot \mathbf{k2}$	$\mathbf{q1} \cdot \mathbf{k3}$	$\mathbf{q1} \cdot \mathbf{k4}$
$\mathbf{q2} \cdot \mathbf{k1}$	$\mathbf{q2} \cdot \mathbf{k2}$	$\mathbf{q2} \cdot \mathbf{k3}$	$\mathbf{q2} \cdot \mathbf{k4}$
$\mathbf{q3} \cdot \mathbf{k1}$	$\mathbf{q3} \cdot \mathbf{k2}$	$\mathbf{q3} \cdot \mathbf{k3}$	$\mathbf{q3} \cdot \mathbf{k4}$
$\mathbf{q4} \cdot \mathbf{k1}$	$\mathbf{q4} \cdot \mathbf{k2}$	$\mathbf{q4} \cdot \mathbf{k3}$	$\mathbf{q4} \cdot \mathbf{k4}$

N

N

Bidirectional vs. Unidirectional Self-Attention

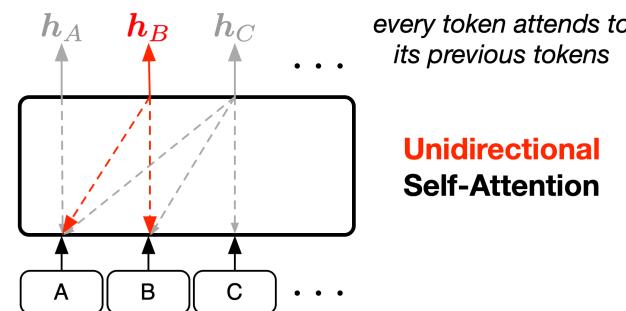
- Self-attention can capture different context dependencies
- **Bidirectional self-attention:**
 - Each position to attend to all other positions in the input sequence
 - Transformers with bidirectional self-attention are called Transformer **encoders** (e.g., BERT)
 - Use case: natural language understanding (NLU) where the entire input is available at once, such as text classification & named entity recognition



Bidirectional vs. Unidirectional Self-Attention

- Self-attention can capture different context dependencies
- **Unidirectional (or causal) self-attention:**
 - Each position can only attend to earlier positions in the sequence (including itself).
 - Transformers with unidirectional self-attention are called Transformer **decoders** (e.g., GPT)
 - Use case: natural language generation (NLG) where the model generates output sequentially

upper-triangle portion set to -inf



N

q1·k1	-∞	-∞	-∞
q2·k1	q2·k2	-∞	-∞
q3·k1	q3·k2	q3·k3	-∞
q4·k1	q4·k2	q4·k3	q4·k4

Position Encoding

- Motivation: inject positional information to input vectors

$$\mathbf{q}_i = \mathbf{x}_i \mathbf{W}^Q \quad \mathbf{k}_i = \mathbf{x}_i \mathbf{W}^K \quad \mathbf{v}_i = \mathbf{x}_i \mathbf{W}^V \in \mathbb{R}^d$$

$$\mathbf{a}_i = \text{Softmax} \left(\frac{\mathbf{q}_i \cdot \mathbf{k}_j}{\sqrt{d}} \right) \cdot \mathbf{v}_j \quad \text{When } \mathbf{x} \text{ is word embedding, } \mathbf{q} \text{ and } \mathbf{k} \text{ do not have positional information!}$$

- How to know the word positions in the sequence? Use position encoding!

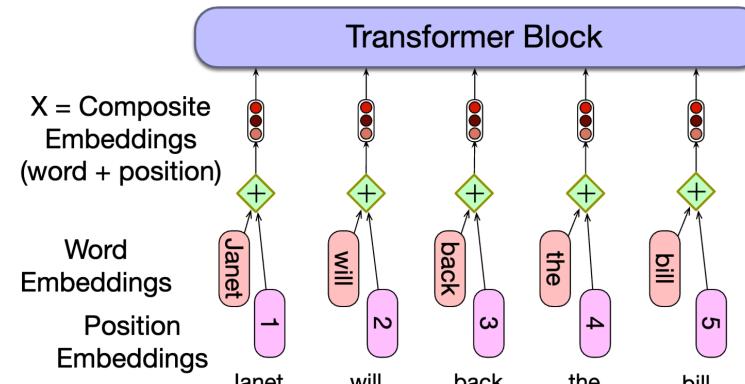


Figure source: <https://web.stanford.edu/~jurafsky/slp3/8.pdf>

Position Encoding Methods

- Absolute position encoding (the original Transformer paper)
 - Learn position embeddings for each position
 - Not generalize well to sequences longer than those seen in training
- Relative position encoding ([Self-Attention with Relative Position Representations](#))
 - Encode the relative distance between words rather than their absolute positions
 - Generalize better to sequences of different lengths
- Rotary position embedding ([RoFormer: Enhanced Transformer with Rotary Position Embedding](#))
 - Apply a rotation matrix to the word embeddings based on their positions
 - Incorporate both absolute and relative positions
 - Generalize effectively to longer sequences
 - Widely-used in latest LLMs



Thank You!

Yu Meng
University of Virginia
yumeng5@virginia.edu